

✓ Hebbian Learning

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 # Display settings
5 %matplotlib inline
6 np.set_printoptions(precision=3, suppress=True)
7
8 seed = 42
9 np.random.seed(seed)
10 # For reproducibility
11 rng = np.random.default_rng(seed)
12
13 # hints: rng.multivariate_normal gives you 2-D gaussians.
14 # .quiver let's you plot actual vector arrows.
```

✓ Exercise 1 • Sampling inputs & visualisation

For each of the following zero-mean 2-D Gaussian input distributions

$$\Sigma_1 = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}, \quad \Sigma_2 = \begin{pmatrix} 1 & 0.4 \\ 0.4 & 1 \end{pmatrix}, \quad \Sigma_3 = \begin{pmatrix} 1 & 0.9 \\ 0.9 & 1 \end{pmatrix}$$

1. Draw **500 samples** and show them in a single scatter plot (use different colours/markers).
2. Draw one random **weight vector** $\mathbf{w} \in \mathbb{R}^2$ from a standard normal distribution and add it to the scatter plot.
3. In a *second* panel visualise the linear activation

$$v = \mathbf{w} \cdot \mathbf{u}$$

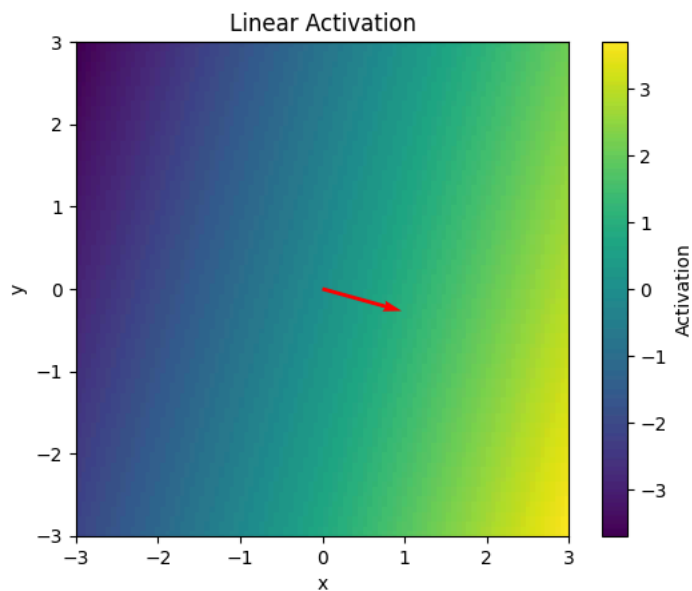


on a 2-D grid covering $[-3, 3] \times [-3, 3]$ using a heatmap, together with the weight vector.

```

1 # hints: rng.multivariate_normal gives you 2-D gaussians.
2 # quiver lets you plot actual vector arrows.
3
4 num_samples = 500
5 mean = [0, 0]
6 covs = [[[1, 0], [0, 1]], [[1, 0.4], [0.4, 1]], [[1, 0.9], [0.9, 1]]]
7
8 # create the 3 datasets
9 datasets = []
10 for distribution in covs:
11     datasets.append(rng.multivariate_normal(mean, distribution, num_samples))
12
13 # plot all 3 datasets
14 for i, data in enumerate(datasets):
15     plt.scatter(data[:, 0], data[:, 1], label=f"Dataset {i + 1}")
16
17
18 # draw and plot weight vector
19 weight = np.random.normal(size=2)
20 weight = weight / np.linalg.norm(weight)
21
22 print(weight)
23 plt.quiver(0, 0, *weight, color="red", angles="xy", scale_units="xy", scale=1)
24 plt.legend()
25 plt.show()
26
27
28 # create activations
29 activations = []
30 for dataset in datasets:
31     activations.append(weight * dataset)
32
33 # plot all 3 activations
34 for i, data in enumerate(activations):
35     plt.scatter(data[:, 0], data[:, 1], label=f"Dataset {i + 1}")
36
37 plt.legend()
38 # create grid for heatmap
39 x = np.linspace(-3, 3, 100)
40 y = np.linspace(-3, 3, 100)
```

```
40     y = np.linspace(0, 5, 100)
41     X, Y = np.meshgrid(x, y)
42     grid_points = np.column_stack((X.ravel(), Y.ravel()))
43
44     # calculate activation for each point
45     activations = np.dot(grid_points, weight)
46     activations = activations.reshape(X.shape)
47
48     # plot heatmap
49     plt.figure()
50     plt.imshow(activations, extent=[-3, 3, -3, 3], origin="lower", aspect="equal")
51     plt.colorbar(label="Activation")
52
53     # plot weight vector on heatmap
54     plt.quiver(0, 0, *weight, color="red", angles="xy", scale_units="xy", scale=1)
55     plt.title("Linear Activation")
56     plt.xlabel("x")
57     plt.ylabel("y")
58     plt.show()
59
```

[0.905 -0.200]



✓ Exercise 2 • Dynamics of Hebbian plasticity

2a – Simple Hebb rule

Implement Hebbian learning

$$\begin{aligned}\Delta \mathbf{w}_n &= v \mathbf{u} \\ v_n &= \mathbf{w}_n^\top \mathbf{u} \\ \mathbf{w}_{n+1} &= \mathbf{w} + \eta \Delta \mathbf{w}_n\end{aligned}$$

- Use a learning rate $\eta = 10^{-3}$.
- Start from a different random initial vector \mathbf{w}_0 .
- Update once per randomly drawn input sample (2000 steps should suffice).
- **Plot** the trajectory of the weight vector on top of the corresponding input scatter for **each** of the three input distributions (use a marker for every 10 updates to keep the plot readable).

Describe qualitatively what you observe.


```

1 def hebb_update(w, u, lr):
2     v = w * u
3     grad = v * u
4     return w + lr * grad
5
6
7 # you can use this function for other learning rules by changing update_fn. cov determines inputs.
8 def run_learning(update_fn, cov, lr=1e-3, steps=2000):
9     # initialize weights, draw a data point, update weights. save weights to a vector.
10    weights = np.random.normal(size=2)
11    weights = weights / np.linalg.norm(weights)
12    traj = []
13    for i in range(steps):
14        data = rng.multivariate_normal([0, 0], cov)
15        weights = update_fn(weights, data, lr)
16        if (i + 1) % 10 == 0:
17            traj.append(weights.copy())
18    return traj
19    return traj
20
21
22 # it's helpful to standardize your plotting for the rest of the experiments.
23 # (ax=None lets you pass in subfigure axis)
24 def plot_trajectory(traj, data, title, ax=None):
25     if ax is None:
26         fig, ax = plt.subplots()
27     ax.scatter(data[:, 0], data[:, 1], alpha=0.2, c="gray", label="Data")
28
29     # Plot trajectory as connected line segments with larger dots
30    traj = np.array(traj)
31    ax.plot(traj[:, 0], traj[:, 1], "b-", alpha=0.5, label="Trajectory", linewidth=2)
32
33    # Plot from the origin to the last trajectory point
34    last_traj = traj[-1]
35    ax.quiver(0, 0, last_traj[0], last_traj[1], color="red", scale=10, label="Last Trajectory")
36    ax.set_title(title)
37    ax.set_xlim(-5, 5)
38    ax.set_ylim(-5, 5)
39    ax.legend()
40    plt.show()
41
42
43 # run the experiment for each cov. matrix and produce the plot for each condition.
44 for i in range(3):
45     traj = run_learning(hebb_update, covs[i])
46     plot_trajectory(traj, datasets[i], f"Weight trajectory for distribution {i + 1}")
47     print(len(traj))
48

```




Weight trajectory for distribution 1




200

Weight trajectory for distribution 2



200

Weight trajectory for distribution 3



200

✓ 2b – Catch your weight vectors!

Repeat part (a) but clip the weight vector to a fixed maximum length

$$\|\mathbf{w}\| \leq w_{\max}^2.$$

```
1 # implement new plasticity update
2 def hebb_clip_update(w, u, lr, w_max=1):
3     v = np.dot(np.transpose(w), u)
4     grad = np.dot(v, u)
5     updated_w = w + lr * grad
6     return updated_w / np.maximum(1, np.linalg.norm(updated_w) / w_max)
7
8
9 # repeat experiment (use run_learning and plot_trajectory)
10
11 # run the experiment for each cov. matrix and produce the plot for each condition.
12 for i in range(3):
13     weight_history = run_learning(hebb_clip_update, covs[i])
14     plot_trajectory(weight_history, datasets[i], " sample title")
```




sample title

2 c — Oja's rule

Implement Oja's *normalising* rulewith $\alpha = 1$.Repeat the experiment for each distribution and verify that the norm of (\mathbf{w}) converges.

Show analytically that at equilibrium

(Hint: take the derivative of $\|\mathbf{w}\|^2$ and set it to zero).


```

1 # implement new plasticity update
2 def oja_update(w, u, lr, alpha=1.0):
3     v = w @ u
4     return w + lr * (v * u - alpha * v**2 * w)
5
6
7 final_norms = []
8
9 for i in range(3):
10     traj = run_learning(oja_update, covs[i])
11     plot_trajectory(traj, datasets[i], "sample title")
12     final_norms.append(np.linalg.norm(traj[-1]))
13
14 # repeat experiment (use run_learning and plot_trajectory)
15 print("Final weight norms()", np.round(final_norms, 3))

```



sample title



sample title

